

Start coding or [generate](#) with AI.

```
import tensorflow as tf
print("GPU Available:", tf.config.list_physical_devices('GPU'))
```

GPU Available: []

```
!nvidia-smi
```

/bin/bash: line 1: nvidia-smi: command not found

```
with tf.device('/GPU:0'):
    print("GPU Available:", tf.config.list_physical_devices('GPU'))
```

GPU Available: []

```
import tensorflow as tf
from tensorflow import keras
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
# import zipfile
# zip_ref = zipfile.ZipFile('/content/drive/MyDrive/Level 6/Artificial_Intelligence/Week5/FruitinAmazon.zip', 'r')
```

```
# zip_ref.extractall('/content')
# zip_ref.close()

# data.data.districts
import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten, BatchNormalization, Dropout

train_dir = '/content/drive/MyDrive/Level 6/AI & ML/w6/FruitinAmazon/train'


img_height = 224 # Example image height
img_width = 224 # Example image width
batch_size = 32
validation_split = 0.2

rescale = tf.keras.layers.Rescaling(1./255) # Normalize pixel values to [0, 1]

train_ds = tf.keras.preprocessing.image_dataset_from_directory(
    train_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    interpolation='nearest',
    batch_size=batch_size,
    shuffle=True,
    validation_split=validation_split,
    subset='training',
    seed=123
)
train_ds = train_ds.map(lambda x, y: (rescale(x), y))
```


➡ Found 90 files belonging to 6 classes.
Using 72 files for training.

```
# Create validation dataset with normalization
val_ds = tf.keras.preprocessing.image_dataset_from_directory(
train_dir,
labels='inferred',
label_mode='int',
image_size=(img_height, img_width),
interpolation='nearest',
batch_size=batch_size,
shuffle=False,
validation_split=validation_split,
subset='validation',
seed=123
)
val_ds = val_ds.map(lambda x, y: (rescale(x), y))
```

 Found 90 files belonging to 6 classes.
Using 18 files for validation.

```
from tensorflow.keras.applications import VGG16
base_model = VGG16(weights='imagenet', include_top=False, input_shape=(224, 224, 3))

base_model.summary()
```

 Model: "vgg16"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool1 (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool1 (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool1 (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool1 (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool1 (MaxPooling2D)	(None, 7, 7, 512)	0

Total params: 14,714,688 (56.13 MB)

```
for layer in base_model.layers:
    layer.trainable = False
```

```
model = Sequential()
model.add(base_model)
model.add(Flatten())
model.add(Dense(1024, activation='relu'))
model.add(Dense(6, activation='softmax'))
```

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 7, 7, 512)	14,714,688
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 1024)	25,691,136
dense_1 (Dense)	(None, 6)	6,150

Total params: 40,411,974 (154.16 MB)

Trainable params: 25,697,286 (98.03 MB)

Non-trainable params: 14,714,688 (56.13 MB)

```
model.compile(optimizer='adam', loss='sparse_categorical_crossentropy', metrics=['accuracy'])
```

```

history = None
with tf.device('/GPU:0'):
    print("GPU AVIALBLE")
    history = model.fit(train_ds, validation_data=val_ds, epochs=10)

```

```

GPU AVIALBLE
Epoch 1/10
3/3 ————— 60s 20s/step - accuracy: 0.6957 - loss: 2.4278 - val_accuracy: 0.5000 - val_loss: 4.6930
Epoch 2/10
3/3 ————— 59s 19s/step - accuracy: 0.7383 - loss: 1.2681 - val_accuracy: 0.7778 - val_loss: 1.0839
Epoch 3/10
3/3 ————— 81s 18s/step - accuracy: 0.9332 - loss: 0.2374 - val_accuracy: 0.9444 - val_loss: 0.3971
Epoch 4/10
3/3 ————— 82s 19s/step - accuracy: 0.8641 - loss: 0.5385 - val_accuracy: 0.8889 - val_loss: 0.5398
Epoch 5/10
3/3 ————— 91s 23s/step - accuracy: 0.9705 - loss: 0.0740 - val_accuracy: 0.8889 - val_loss: 0.5773
Epoch 6/10
3/3 ————— 71s 18s/step - accuracy: 0.9891 - loss: 0.0139 - val_accuracy: 0.8889 - val_loss: 0.6613
Epoch 7/10
3/3 ————— 58s 18s/step - accuracy: 1.0000 - loss: 0.0112 - val_accuracy: 0.8889 - val_loss: 0.6499
Epoch 8/10
3/3 ————— 69s 25s/step - accuracy: 1.0000 - loss: 0.0060 - val_accuracy: 0.8889 - val_loss: 0.5778
Epoch 9/10
3/3 ————— 69s 18s/step - accuracy: 1.0000 - loss: 9.5737e-04 - val_accuracy: 0.8889 - val_loss: 0.5238
Epoch 10/10
3/3 ————— 58s 18s/step - accuracy: 1.0000 - loss: 4.7257e-04 - val_accuracy: 0.8889 - val_loss: 0.4811

```

```

test_loss, test_acc = model.evaluate(val_ds)
print(f"Validation Accuracy: {test_acc:.2f}")

```

```

1/1 ————— 13s 13s/step - accuracy: 0.8889 - loss: 0.4811
Validation Accuracy: 0.89



```

```
#plot for test data
# Plot training history
import matplotlib.pyplot as plt
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.legend()

plt.show()
```



```
test_dir = "/content/drive/MyDrive/Level 6/AI & ML/w6/FruitinAmazon/test"
test_ds = tf.keras.preprocessing.image_dataset_from_directory(
    test_dir,
    labels='inferred',
    label_mode='int',
    image_size=(img_height, img_width),
    batch_size=batch_size,
    shuffle=False,
    interpolation='nearest',
    seed=123
)
```

```
test_ds = test_ds.map(lambda x, y: (rescale(x), y))
```


➞ Found 30 files belonging to 6 classes.

```
test_loss, test_accuracy = model.evaluate(test_ds)
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test Loss: {test_loss:.4f}")
```

➞ 1/1 ————— 27s 27s/step - accuracy: 0.5667 - loss: 3.0391
Test Accuracy: 0.5667
Test Loss: 3.0391

```
import tensorflow
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Flatten
from keras.applications.vgg16 import VGG16
from tensorflow.keras.optimizers import RMSprop
```

```
base_model = VGG16(
    weights='imagenet',
    include_top = False,
    input_shape=(224,224,3)
)
```

```
base_model.trainable = True
```


```
set_trainable = False
```

```
for layer in base_model.layers:
    if layer.name == 'block5_conv1':
        set_trainable = True
    if set_trainable:
        layer.trainable = True
    else:
        layer.trainable = False
```

```
for layer in base_model.layers:  
    print(layer.name,layer.trainable)
```

```
⇒ input_layer_2 False  
   block1_conv1 False  
   block1_conv2 False  
   block1_pool False  
   block2_conv1 False  
   block2_conv2 False  
   block2_pool False  
   block3_conv1 False  
   block3_conv2 False  
   block3_conv3 False  
   block3_pool False  
   block4_conv1 False  
   block4_conv2 False  
   block4_conv3 False  
   block4_pool False  
   block5_conv1 True  
   block5_conv2 True  
   block5_conv3 True  
   block5_pool True
```

```
base_model.summary()
```

 Model: "vgg16"

Layer (type)	Output Shape	Param #
input_layer_2 (InputLayer)	(None, 224, 224, 3)	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1,792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36,928
block1_pool1 (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73,856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147,584
block2_pool1 (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295,168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590,080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590,080
block3_pool1 (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1,180,160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2,359,808
block4_pool1 (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2,359,808
block5_pool1 (MaxPooling2D)	(None, 7, 7, 512)	0

Total params: 14,714,688 (56.13 MB)

```
model = Sequential()

model.add(base_model)
model.add(Flatten())
model.add(Dense(256,activation='relu'))
model.add(Dense(6,activation='softmax'))

model.compile(optimizer=keras.optimizers.RMSprop(learning_rate=1e-5), loss='sparse_categorical_crossentropy', metrics=['accuracy'])



test_loss, test_acc = model.evaluate(val_ds)
print(f"Validation Accuracy: {test_acc:.2f}")
```

➡ 1/1 ————— 13s 13s/step - accuracy: 0.1667 - loss: 2.8538
Validation Accuracy: 0.17

```
#plot for test data
# Plot training history
import matplotlib.pyplot as plt
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Training Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Loss')
plt.legend()
```



```
test_dir = "/content/drive/MyDrive/Level 6/AI & ML/w6/FruitinAmazon/test"
```

```
test_ds = tf.keras.preprocessing.image_dataset_from_directory(
```